

```

      .
|;;|                                     |;;|
|[ ]-----[ ]|
|;;|               -----             |;;| | | |
|;;|           -----                 |;;|
|;;|   | COMP6447 |                   |;;|
|;;|       -----                     |;;|
|;;| learning c                       |;;|
|;;|_____                           |;;|
|;;;;;;;;;;;;;|                      |;;;;;;;;;;;;;|
|;;;;;;;;;|_____|                    |;;;;;;;;;|
|;;;;;;;;;|    | |;;;|                |;;;;;;;;;|
|;;;;;;;;;|    | |;;;|                |;;;;;;;;;|
|;;;;;;;;;|    | |;;;|                |;;;;;;;;;|
|;;;;;;;;;|    | |;;;|                |;;;;;;;;;|
|;;;;;;;;;|    | |;;;|                |;;;;;;;;;|
|;;;;;;;;;|    |___|                  |;;;;;;;;;|
\_____|_____|\_____|_____||
~~~~~^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ~~~~~

```

Source code auditing

And **fuzzzzzzing**

Why do we read the source

1. Because it's the best way to find bugs
2. Because it's easier than reversing assembly
3. A lot of people prefer to fuzz and manipulate input
 - a. They won't find the bugs you will

Writing exploits is **easier** if you **understand** the **program**

Read the Source luke

Several classes of problems

- Bad api usage
- Logic bugs
- Integer overflows
- Type conversions / confusions
- Array out of bounds access
- Incorrect use of operators (sizeof / bit shifts / etc)
- Not understanding pointer arithmetic
- Race conditions

Let's talk about **some** of these

Bad api usage

- Using apis in **stdlib** assume the **programmer** has a **brain**
 - Big mistake
- Not checking sizes
- Classic format string

```
Int main(int argc, char* argv[]) {  
  
    printf(argv[1]);  
  
}
```

Ordering parameters wrong...

strcpy?

```
Int main(int argc, char* argv[]) {  
    Char s[4];  
    strcpy(s, argv[1]);  
}
```

More common than gets, less obvious to the developer why its bad

But im a good dev adam. I use **strncpy**

```
char *strncpy(char *dest, const char *src, size_t n);
```

- Doesn't copy NULL byte if the size doesn't allow it

Hmm...

Whats the problem here

```
int main(int argc, char* argv[]) {  
    Char important[16] = "ABCDEFGHIKLMNO";  
  
    Char b[16];  
  
    strncpy(b, argv[1], 16);  
  
    printf("%s\n", b);  
}
```


Do you think developers are good?

- The way C is compiled / typed, sometimes using an API a wrong way is a valid way to use it
- Look at `memset`
 - `memset(void* s, int c, size_t n);`
- Fills (**s**) with **n** bytes of **c**

Seems reasonable right?

Does this zero out memory?

- **`memset(s, 100, 0)`**

Integer overflows and underflows...

```
Unsigned int a;
```

```
inr a = INT_MAX
```

```
a += 1
```

What's a?

```
unsigned int a = 0
```

```
a--;
```

What's A?

```
u_char *   make_table(unsigned int width, unsigned int height,  
u_char *init_row)  
{  
    unsigned int n;  
    int i;  
    u_char *buf;  
    n = width * height;  
    buf = (char *)malloc(n);  
  
    if (!buf) return (NULL);  
  
    for (i=0; i< height; i++)  
        memcpy(&buf[i*width], init_row, width);  
    return buf;  
}
```

What happens if you multiply 2 big numbers

What happens when you multiply two massive numbers?

$0x400 * 0x1000001 = 0x400000400$ overflows to become $0x400$ lol

Easy overflow.. Easy to spot

This is from an open source tool...

```
u_int nresp;  
...  
nresp = packet_get_int();  
if (nresp > 0) {  
    response = xmalloc(nresp * sizeof(char*));  
    for (i = 0; i < nresp; i++)  
        response[i] = packet_get_string(NULL);  
}  
packet_check_eom();
```

What does `sizeof(char*)` return?

length+1 might overflow

```
char *read_data(int sockfd)
{
    char *buf;
    int length = network_get_int(sockfd);
    if(!(buf = (char *)malloc(MAXCHARS)))
        die("malloc: %m");
    if(length < 0 || length + 1 >= MAXCHARS){
        free(buf);
        die("bad length: %d", value);
    }
    if(read(sockfd, buf, length) <= 0){
        free(buf);
        die("read: %m");
    }
    buf[value] = '\\0';
    return buf;
}
```

is this sudo vulnerable?

```
void assume_privs(unsigned short uid)
{
    seteuid(uid);
    setuid(uid);
}

int become_user(int uid)
{
    if (uid == 0)
        die("root isnt allowed");
    assume_privs(uid);
}
```

- Become_user takes an int
- Assume_priv takes a short
- An int like 0x1000 0000 would get truncated to 0x0000
 - The first check passes
 - But they still get root
 - Guess where i pulled this example from lol

What is a race condition?

A race condition is a **generic situation** where **two or more actors** are about to perform an action, and the result depends on the order in which they occur

example

- Checks file doesn't exist
- Does some stuff
- Opens file
- Appends to file
- Saves file

```
#define DELAY 10000
```

```
int main() {
    char *fn = "/tmp/XYZ";
    char buffer[60];
    FILE *fp;

    long int i;
    /* get user input */
    scanf("%50s", buffer);
    if (!access(fn, W_OK)) {
        /* simulating delay */
        for (i = 0; i < DELAY; i++) {
            int a = i^2;
        }

        fp = fopen(fn, "a+");
        fwrite("\n", sizeof(char), 1, fp);
        fwrite(buffer, sizeof(char), strlen(buffer), fp);
        fclose(fp);
    }

    else
        printf("No permission \n");
}
```

example

- Checks file doesn't exist
- Does some stuff
- **Attacker creates symlink it to /etc/shadow**
- Opens file
- Appends to file
- Saves file

```
#define DELAY 10000

int main() {
    char *fn = "/tmp/XYZ";
    char buffer[60];
    FILE *fp;

    long int i;
    /* get user input */
    scanf("%50s", buffer);
    if (!access(fn, W_OK)) {
        /* simulating delay */
        for (i = 0; i < DELAY; i++) {
            int a = i^2;
        }

        fp = fopen(fn, "a+");
        fwrite("\n", sizeof(char), 1, fp);
        fwrite(buffer, sizeof(char), strlen(buffer), fp);
        fclose(fp);
    }

    else
        printf("No permission \n");
}
```

Source auditing is similar to reversing

- Look at the big picture
- Audit functions you know are unsafe
 - Printf, strcpy, gets
- Make a list of interesting functions as you audit
 - Try to understand what a function does by its name first
 - Don't jump between functions too much, you'll confuse yourself
- Bottom up approach
 - Start at user input
 - Move up and see where your data goes
- Read the manual



Fuzzing

- Fuzzing or fuzz testing is an automated software testing technique that involves providing invalid, unexpected, or random data as inputs to a computer program.
- Not always used in a security sense, can be used for unit testing of software.

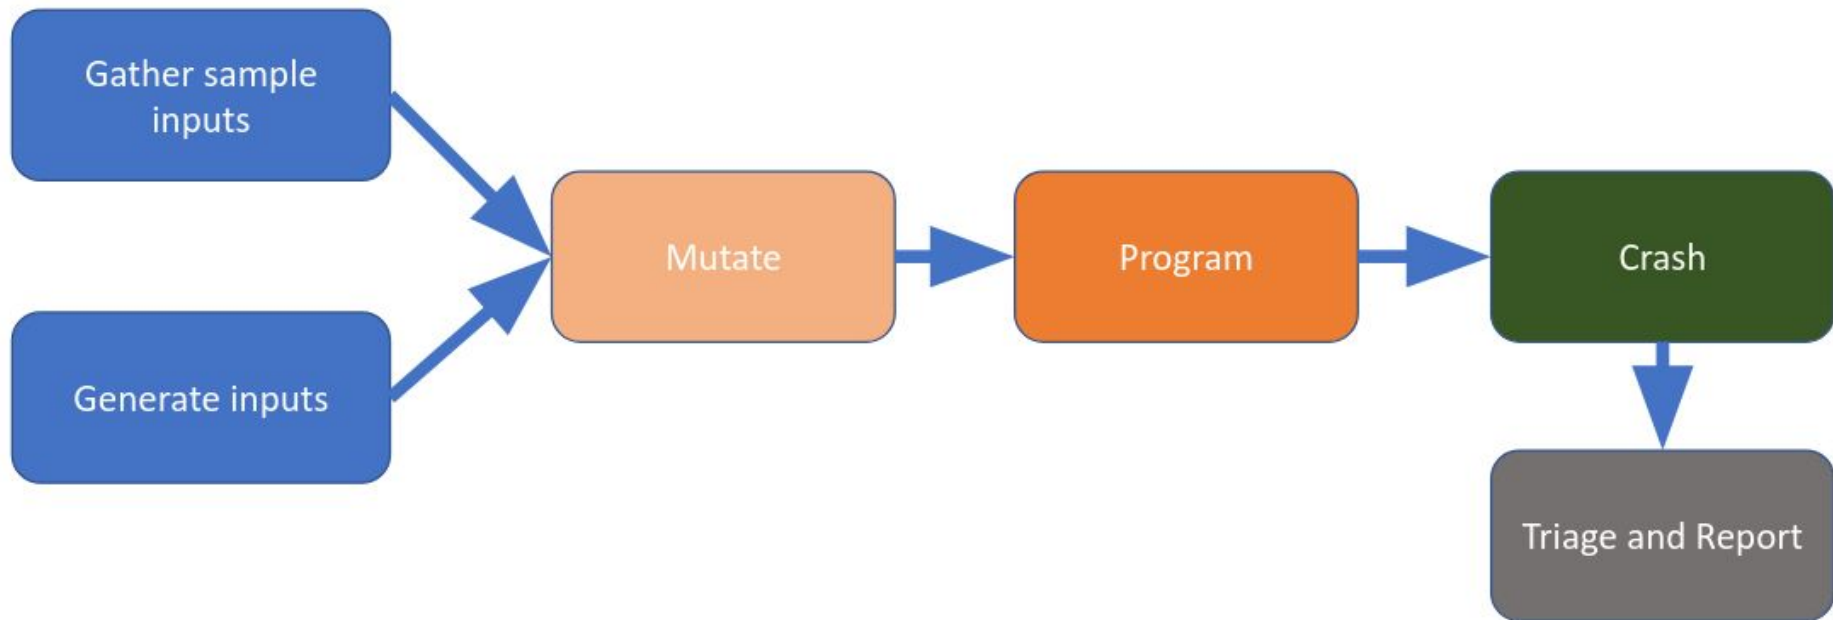
tldr;

- automate the finding of vulnerabilities
 - Generate test cases for program
 - Send them to program
 - Monitor for any changes in the programs behaviour
- Goal is to get the most code **coverage**
 - If you can test as many code **paths** as possible, more likely to find bugs
 - Monitoring code coverage is hard

More on fuzzing

- Fuzzing is often used to test file formats or networking protocols
 - Structured inputs
- Usually a fuzzer takes in a sample input, and mutates it
 - Or takes in **large sets** of inputs, and mutates them
- Goal is to create an exceptional condition
 - Program crashes
 - Hangs (infinite loops?)
 - Unique errors/error codes

Fuzzer cycle



Fuzzing has limitations compared to auditing

- Custom protocols / file formats are hard to know about without prior reversing
- Non crashing bugs are tough to detect
 - Arbitrary file read/write
 - SQL injection
 - Priv escs
- Hard to know what progress you have
 - Sometimes you will get no crashes for a long time.
 - Does this mean no bugs?
 - Is your fuzzer broken?
 - Sometimes you will get **a lot** of crashes
 - Maybe your fuzzer is triggering the same bug over and over
 - Not finding more useful bugs
- Most off the shelf fuzzers require source code for the program
- Generating all possible inputs to a program is not possible

Fuzzing has 2 main parts

- The fuzzer
- The harness

Harness

- The thing you use to watch the program behave
 - Could be a **debugger**
 - Could be **code coverage** stuff built into the binary
 - Can use a **hypervisor** to fuzz an entire system
 - Could just look at the **return code** after executing the program
- Stuff to look for
 - Crashes
 - Weird program states
 - Code coverage
 - Error messages
 - Information leaks?
- tldr: monitors for crashes, records what inputs caused what crash

How to start out writing your own fuzzer

Start by just running the program

- Write a tool that takes some data, mutates it, and saves the mutated file
- Run the program to parse the mutated file
- Did the program crash?
 - Attach gdb or enable core dumps :D
 - Or look at return codes
- This has some big negatives
 - Can't really scale, if the program relies on one file
 - Program startup is slow

Open source fuzzers

AFL

- One of the most used fuzzers
 - Libfuzzer is also good
- **Relies on source being available***
 - Looks at gaining most coverage
- Takes in sample tests (corpus based)
 - Mutates them
 - Looks at output
 - Keeps mutated files that reveals new code paths
 - Repeatedly mutate the file using a balanced and well-researched variety of traditional fuzzing strategies
- **Coverage guided fuzzing**

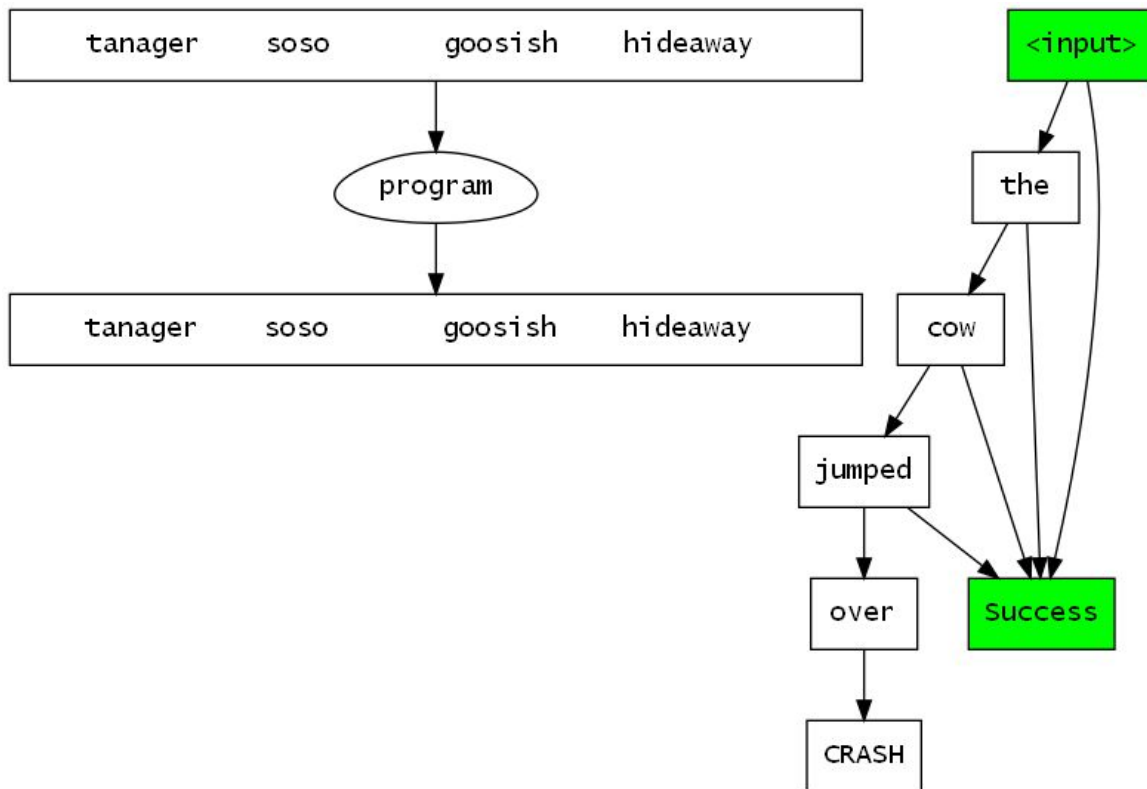
american fuzzy lop 1.18b (png2_extras)

process timing		overall results
run time : 4 days, 22 hrs, 17 min, 16 sec		cycles done : 48
last new path : 0 days, 1 hrs, 43 min, 42 sec		total paths : 3804
last uniq crash : none seen yet		uniq crashes : 0
last uniq hang : 0 days, 0 hrs, 38 min, 38 sec		uniq hangs : 75
cycle progress	map coverage	
now processing : 1298 (34.12%)	map density : 4169 (6.36%)	
paths timed out : 0 (0.00%)	count coverage : 4.82 bits/tuple	
stage progress	findings in depth	
now trying : splice 16	favored paths : 458 (12.04%)	
stage execs : 1050/1500 (70.00%)	new edges on : 710 (18.66%)	
total execs : 961M	total crashes : 0 (0 unique)	
exec speed : 1423/sec	total hangs : 459 (75 unique)	
fuzzing strategy yields	path geometry	
bit flips : 1/423k, 0/423k, 0/423k	levels : 21	
byte flips : 0/52.9k, 0/52.9k, 0/52.9k	pending : 2	
arithmetics : 1/2.95M, 0/1.07M, 0/438k	pend fav : 0	
known ints : 0/316k, 0/1.81M, 0/2.47M	own finds : 17	
dictionary : 0/1.37M, 0/1.38M, 0/2.64M	imported : n/a	
havoc : 10/316M, 5/628M	variable : 0	
trim : 4537 B/1.10M (0.14% gain)		

[cpu:124%]

Coverage based fuzzing

Follow code sections



Finding more important bugs

Not every bug causes a crash

- Small Overflows might not crash
- Integer overflows/underflows might not crash
- Invalid Auth success might not crash

Add hooks into the program at these points, crash if condition is true.

This way you can find when these things happen

Address Sanitizer is one way of doing this, it crashes a program on invalid memory writes/reads

More open source

- Libfuzzer is another popular one
- Part of LLVM, designed to be built with your program
- Really fast because everything is done in memory (no files read from disk)
- Is corpus based similar to AFL

How to make a better fuzzer?

- Improve performance
 - Moar threads
 - Not too many though, don't kill the machine :D
- Most fuzzers are slow
- Relying on disk + syscalls is slow
 - So don't?
 - Ie: Pass files through memory

Example fuzzer input/out

Sample input:

- GET /index.html HTTP/1.1

Fuzzer output:

- AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
- AAAAAAAAAA....AAAAAAAAAA /index.html HTTP/1.1
- GET %n%n%n%n%n%n HTTP/1.1
- GET /index.html HTTTTTTTTTTP/1.1
- GET /index.html HTTP/0.0
- etcetcetc

Example vulnerable code

```
char buffer[1024]
```

```
strcpy(buffer, argv[1])
```

```
return 0;
```

```
while True:
```

```
    subprocess.call(["./program", "A" * i])
```

```
    if crash:
```

```
        print('we got em')
```

Fuzzing is about **intelligently** generating input...

Fuzzing strategies

```
fuzzing strategy yields
bit flips : 1/423k, 0/423k, 0/423k
byte flips : 0/52.9k, 0/52.9k, 0/52.9k
arithmetics : 1/2.95M, 0/1.07M, 0/438k
known ints : 0/316k, 0/1.81M, 0/2.47M
dictionary : 0/1.37M, 0/1.38M, 0/2.64M
havoc : 10/316M, 5/628M
trim : 4537 B/1.10M (0.14% gain)
```

Mutation based fuzzing - Dumb but easy

Generation based fuzzing - Smart but hard

Mutation based fuzzing

- No understanding of structure of input
- Completely random inputs
- Add completely random anomalies to existing input
- Easy and quick to setup

Generation based fuzzing

- Generate input based on some format / test case
 - Example: input must follow some spec, etc
- Anomalies are added at each spot in the input
 - If json, field names are fuzzed, field values are fuzzed, arrays are fuzzed
 - If xml,
 - Situational, you must know what the data looks like
- Knowledge of protocol gives a lot more insight compared to random fuzzing

```
fuzzing strategy yields
bit flips : 1/423k, 0/423k, 0/423k
byte flips : 0/52.9k, 0/52.9k, 0/52.9k
arithmetics : 1/2.95M, 0/1.07M, 0/438k
known ints : 0/316k, 0/1.81M, 0/2.47M
dictionary : 0/1.37M, 0/1.38M, 0/2.64M
havoc : 10/316M, 5/628M
trim : 4537 B/1.10M (0.14% gain)
```

Example fuzzer demo

- Lets hope this demo of a bitflipper works

What do you do if you keep triggering the same bug..

- Add tests to your fuzzer to ignore it
- Try a different technique (never stick to one technique)
- Given a crash.. How do you find the actual vulnerability and exploit it
- When do you give up on fuzzing

Assignment...

Write a fuzzer in groups of 4.

Your goal is to **write a fuzzer** that **given a sample input and a binary** will **generate a file** that when passed into the **program**, **causes a crash**.

Will be given 10 sample binaries + 10 sample inputs.

Will be tested on **more**.

Assignment is intentionally designed with a wide scope. Be creative...

Example test case

```
$ ls
fuzzer.exe input.txt binary
$ ./fuzzer.exe binary input.txt
Fuzzing this thing...
$ ls
fuzzer.exe binary bad.txt input.txt
$ cat bad.txt | binary
Segmentation Fault
```

Marks (30 max marks)

20 marks - for finding all vulnerabilities in the **10 provided binaries**.

10 marks - for finding all vulnerabilities in the **binaries that we do not provide**.

6 marks - Something awesome . Something **cool** your fuzzer does.

Some hints / assumptions

Writing a fully black box fuzzer is hard. You can assume some things about the input:

- Input will be one of the following text formats:
 - Plaintext (multiline)
 - yaml
 - json
 - xml
 - csv
- Bugs might be in the program or the parser :-)